

Topic 12

Cache vs. hash

資料結構與程式設計
Data Structure and Programming

12/04/2019

1

From $O(\log n)$ to $O(1)$?

- ◆ For set and map, they have good complexity for "insert", "delete" and "find" operations
→ $O(\log n)$
- ◆ However, in set and map, all the data are sorted --
 - Can output the data in ascending/descending order
 - Can get the list of elements with values in certain range
- ◆ What if we don't care about the order, and just want to have fast "insert", "delete" and "find" operations?
 - Can we gain something (complexity) back for not sorting the data?

Think: How do you retrieve memory?

Data Structure and Programming


Prof. Chung-Yang (Ric) Huang

2

2

Hash Table

1. Buckets: the table is composed of B buckets (usually a large number)
2. Each bucket can hold up to S slots of data (usually a smaller number)
3. Given a data d with key x , a hash function $h(x)$ is used to compute the corresponding bucket number



Data Structure and Programming


Prof. Chung-Yang (Ric) Huang

3

3

Hash Table Example

- ◆ Record: (student name, score)
- ◆ Hash table: 26 buckets
- ◆ Hash function = the first character of name



Data Structure and Programming

Prof. Chung-Yang (Ric) Huang

4

4

Complexity Analysis

- ◆ Depending on how the s slots are designed
- ◆ However, the worse case...
 - Insert: $O(1)$
 - Assuming it takes $O(1)$ to compute $h(x)$
 - Delete: $O(s)$ → can they be $O(\log s)$?
 - Find: $O(s)$ → what's the price to pay?
- ◆ Because s is usually a smaller number (e.g. 2) → Very efficient

Hash Table Design Issues

1. Choice of hash function
2. Overflow handling methods
3. Size (number of buckets) of hash table

Hash Functions

- ◆ Convert key x to an integer b as the bucket index ($0 \leq b \leq B - 1$)
 - Generally $O(1)$ complexity
- ◆ Discuss: how's the hash function used in slide #4 (student name, score) example?
 - No good, the first characters of names are usually not evenly distributed among 26 letters
 - Some buckets may become full easily (e.g. 'c'), while some may be empty (e.g. 'x')
- ◆ Ideal: for all possible key values, approximately the same number of keys get mapped into each bucket
→ Uniform Hash Function

Hash Function Methods

- ◆ Summation
 - e.g. Adding the ASCII values of some/all the characters together
- ◆ Shift
 - e.g. Keys are pointers; because pointer values are multiplier of 4 (or 8)
→ $h(x) = (x \gg 2) \dots$
 - e.g. $x.p$, $x.q$ are two pointers;
→ $h(x) = (x.p \gg 3) + (x.q \gg 6);$
- ◆ Division
 - e.g. Divide a prime number
- ◆ Others: folding, mid-square, digit analysis, etc
- ◆ Usually: mixed of the above

Collision and Overflow

- ◆ Collision
 - Two non-identical keys are hashed into the same bucket
 - At most $(s - 1)$ collisions for a buckets
 - Reduced by better hash function
 - ◆ Overflow
 - When a new key is hashed into a full bucket
- For better hash performance, we should try to produce less collisions and prevent overflow

Overflow Handling

- ◆ Overflow may still happen when more and more data are stored into the hash
1. Open addressing
 - Find a non-full bucket to insert the new key
 - Linear probing, quadratic probing, rehashing, (pseudo)random probing, etc
 2. Chaining
 - Use linked list, dynamic array, or other kinds of ADT to make the s extendible

Hash Table Size

- ◆ Hash table size (number of buckets) also affects the occurrence of overflow
 - Too small → Overflow happens
 - Too large → Waste of memory
1. Static hashing
 - Fixed-size hash table
 - Easier to implement; better if the number of possible elements is known in advance
 2. Dynamic hashing
 - Hash table size can grow when necessary
 - When it grows, rehashing is needed

HW#7 -- Hash Implementation (myHashSet.h)

```
template <class Data>
class HashSet
{
public:
    class iterator
    {
        friend class HashSet<Data>;
    };
private:
    size_t          _numBuckets;
    // Each _buckets[i] is a vector<Data>
    vector<Data>*   _buckets;
};
```

Must overload:

- operator () → as the hash function (i.e. $d() \% \text{_numBuckets}$ to return bucket index)
- operator == → to compare if the data are equivalent (or have the same key)

Supported functions in class HashSet

```

iterator begin() const; // Point to the first valid data
iterator end() const; // Past the end

bool empty() const; // return true if no valid data

size_t size() const; // number of valid data

vector<Data>& operator [] (size_t i) { return _buckets[i]; }

const vector<Data>& operator [] (size_t i) const;

void init(size_t b); // initialize Hash with _numBuckets = b

void reset();

```

13

Supported functions in class HashSet

```

// check if d is in the hash...
// if yes, return true; else return false;
bool check(const Data& d) const;

// query if d is in the hash...
// if yes, replace d with the data in the hash and return true;
// else return false;
bool query(Data& d) const;

// update the entry in hash that is equal to d (i.e. == return true)
// if found, update that entry with d and return true;
// else insert d into hash as a new entry and return false;
bool update(const Data& d);

// return true if inserted successfully (i.e. d is not in the hash)
// return false if d is already in the hash ==> will not insert
bool insert(const Data& d);

```

14

Another Hash Implementation Example (myHashMap.h)

```

template <class HashKey, class HashData>
class HashMap
{
    typedef pair<HashKey, HashData> HashNode;
public:
    class iterator
    {
        friend class HashMap<HashKey, HashData>;
    };
private:
    size_t          _numBuckets;
    vector<HashNode>* _buckets;
};

```

15

Another Hash Implementation Example (myHashMap.h)

class HashMap<HashKey, HashData>

vector<HashNode>* _buckets

_numBuckets

All the HashNodes have the same bucketNum(), but distinct HashKeys

vector<HashNode>

HashNode
≡ pair<HashKey, HashData> (k, d)

Overload operator () for HashKey as hash function

Usually a prime number (Why?)

size_t bucketNum(const HashKey& k) const { return k() % _numBuckets; }

16

Class HashKey

- ◆ To use Hash ADT, you should define your own HashKey class.
- It should at least overload the "()" and "==" operators.

```
class HashKey // Of course, name your own HashKey
{
public:
    HashKey(); // define your own constructor
    size_t operator() () const; // acted as "hash
                                // function"
    bool operator == (const HashKey& k);
                                // to compare the HashKey

private:
    // Define your own data members
};
```

Example of using class Hash

- ◆ Locating an address
- ◆ `typedef string Address;`
`typedef pair<float, float> Location;`
`class AddressKey {`
`public:`
 `size_t operator() () const { ...; }`
 `bool operator == (const Address& a) {`
 `return (_addr == a._addr); }`
`private:`
 `Address _addr;`
`};`

class Hash::iterator

```
class iterator
{
    friend class HashMap<HashKey, HashData>;
private:
    // Define your own data members!!
};
◆ Purpose: to go through the "valid" HashNodes in the Hash
◆ To use:
    HashMap<HashKey, HashData> hh;
    ... // insert data
    HashMap<HashKey, HashData>::iterator hi = hh.begin();
    for (; hi != hh.end(); ++hi)
        cout << (*hi).first << " → " << (*hi).second
               << endl;
```

Supported functions in class HashMap

```
iterator begin() const; // Point to the first valid data
iterator end() const; // Pass the end
bool empty() const; // return true if no valid data
size_t size() const; // number of valid data
vector<HashNode>& operator [] (size_t i) { return _buckets[i]; }
const vector<HashNode>& operator [] (size_t i) const;
void init(size_t b); // initialize Hash with _numBuckets = b
void reset();
bool check(const HashKey& k, HashData& n);
bool insert(const HashKey& k, const HashData& d);
bool replaceInsert(const HashKey& k, const HashData& d);
void forceInsert(const HashKey& k, HashData d);
```

Hash Classes in STL

- ◆ Since C++-11, STL supports and standardizes the following associative containers that implement Hash

1. unordered_set
2. unordered_multiset
3. unordered_map
4. unordered_multimap

[Ref] Deprecated, non-standard STL implementations:

1. hash_set
2. hash_multiset
3. hash_map
4. hash_multimap

class unordered_set in STL

- ◆ unordered_set<Key[, Hash, Pred, Alloc]>
 - class Key: element type
 - class Hash: hash function (optional; default = hash<Key>)
 - class Pred: equality checking for class Key (optional; default = equal_to<Key>)
 - class Alloc: used for internal memory management (optional; default = alloc)

Member Functions of class unordered_set

1. iterator begin(); const_iterator cbegin() const;
iterator begin(n); // get the first element in bucket[n]
// see also cbegin(n), end(), end(n), cend(), cend(n)
2. size_type bucket_count() const; size_type bucket(k) const;
3. pair<iterator, bool> insert(const value_type& k);
void insert(InputIterator f, InputIterator l);
4. iterator erase(const_iterator pos); // to the next iterator
size_type erase(const key_type& k); // #elms erased
iterator erase(const_iterator first, const_iterator last);
5. iterator find(k) const; size_type count(k) const;
6. float load_factor() const; // = size() / bucket_count()
float max_load_factor() const; void max_load_factor(l);
void rehash(size_type n);

Other Hash Classes in STL

- ◆ class unordered_multiset
 - Similar to unordered_set, but allow elements with identical keys
- ◆ class unordered_map
 - unordered_map<Key, Data[, Hash, Pred, Alloc]>
- ◆ class unordered_multimap
 - Similar to unordered_map, but allow elements with identical keys

Cache in Programming

- ◆ Structurally similar to hash, however
 - Usually smaller number of buckets
 - Each bucket contains exactly 1 element
 - When collision happens, the old data is overwritten by the new one
 - Easier to implement than hash
- ◆ Usually used as computational cache
 - (input parameters) → computational results
- ◆ There is no STL implement yet

Cache Implementation in util/myHashMap.h

```
template <class CacheKey, class CacheData>
class Cache
{
    typedef pair<CacheKey, CacheData> CacheNode;
public:
    // No need to implement class iterator (why?)
    void init(size_t s);
    void reset();
    size_t size() const;
    CacheNode& operator [] (size_t i);
    const CacheNode& operator [] (size_t i) const;
    bool read(const CacheKey& k, CacheData& d) const;
    void write(const CacheKey& k, const CacheData& d);
private:
    size_t      _size;
    CacheNode*  _cache; // new CacheNode[_size]
};
```

Example of using class Cache (BDD project)

- ◆ Computed table
$$\text{ITE}(F, G, H) = F * G + \overline{F} * H$$

F, G, H: BddNode (contains a size_t data)

 - Requires expensive recursive calls to compute ITE() functions ($O(|H|*|G|*|H|)$)
 - The computed table (cache) is to record the result (as CacheData) with respect to the ITE parameters (as CacheKey)
 - So next time when the same ITE(F, G, H) computation is required, we can immediately look up the cached result